

# Reversible JPEG Coefficient Swapping: Reducing Restoration-Information Cost with ARPT and Confidence-Guided Carrier Selection

Shao-An Chien

Preliminary manuscript — not peer reviewed

Date: 2026-10-05

Department of Computer Science and Information Engineering, Ming Chuan University  
Email: 11360483@mcu.edu.tw, rex99075@gmail.com

## Abstract

Reversible watermarking attaches authorization information to an image for copyright protection and restores the original image after extraction for subsequent use. JPEG is a common format for image sharing. Existing AC-pair swapping methods can provide high watermark robustness, but restoring the original image requires recording prediction errors; reducing these errors therefore lowers the restoration-information cost.

This paper proposes an Antisymmetric Residual-Pairing Transformer (ARPT) that compares two candidate coefficient-pair arrangements to predict the swap state. Carriers are selected according to prediction confidence to reduce the number of errors that must be recorded. The remaining errors are stored as a lossless correction map in the marked JPEG, allowing the receiver to restore the original image.

Experiments show that ARPT achieves approximately 94.7% swap-state classification accuracy. Confidence-guided carrier selection reduces the error rate among selected carriers from 5.45% to 0.66%, lowers the restoration-information cost, and permits exact restoration of the original quantized coefficients and decoded pixels.

**Keywords:** reversible watermarking; JPEG; ARPT; confidence-guided carrier selection; correction map; overhead-adjusted payload

## 1. Introduction

Digital watermarking embeds copyright labels or authorization information in an image so that the information is retained and transmitted with the image. For copyright protection, a watermark should resist malicious removal so that the associated information can still be read and verified after the image has circulated. Watermark embedding also changes the image, however. If the original data are needed later, the image must be restored to its pre-embedding state after message extraction. This is the purpose of reversible watermarking. JPEG is a common format for image storage and sharing, and this study addresses reversible watermarking in JPEG images.

Recent research on reversible data hiding in JPEG has largely examined how to balance image distortion and file size while carrying a message. Weng et al. (2023) [1], for example, represented two AC coefficients in a block as two-dimensional coordinates and generated modification rules adaptively according to block texture to reduce distortion caused by large coefficient changes. Li et al. (2024) [2] proposed progressive modification, activating coefficients at selected frequencies according to the required message size to avoid unnecessary shifts in coefficients that carry no message and thus limit file expansion. These histogram-modification methods are highly sensitive to changes in coefficient values, however, which makes it difficult to meet the interference-resistance requirements of copyright protection.

To address these requirements, Ou and Chen (2016) [3] studied watermarking for stereo-pair images and proposed embedding watermarks through AC-pair swapping. Their method constructs a feature bitmap from differences between the left and right images, then embeds the watermark in the left image by swapping the order of designated AC coefficients. During verification, the receiver reads the watermark and compares it with features recomputed from the images to determine whether the copyright marking is consistent. Their experiments showed that the watermark could still be read and verified under the malicious attacks they tested, demonstrating the robustness of that method.

To restore an image after watermark extraction, Ou (2024) [4] further investigated the reversibility of AC-pair swapping. The method represents a bit by the relative order of a pair of quantized AC coefficients, so the receiver can read the message directly from that order. For restoration, spatial-domain IDCT rounding residuals are used to determine whether the coefficients were swapped. The locations of incorrectly classified blocks are recorded in the watermark so that those blocks can be corrected by swapping the pair during restoration. This design makes complete restoration of the original image possible while retaining the robustness of coefficient-swapping watermarking.

The restoration procedure nevertheless requires recording every prediction error, and this information directly occupies space that would otherwise be available for message embedding. Ou (2024) [4] reported that restoration information accounted for 10%–20% of the available embedding bits in most test images, with some images approaching the reserved upper limit of 25%. In an architecture that embeds only one bit per block, this cost greatly reduces effective payload efficiency. Improving swap-state classification accuracy can reduce the number of errors that must be stored and hence the restoration-information cost. This study measures that benefit using overhead-adjusted payload.

We propose ARPT to improve swap-state classification. Two candidates are constructed for each block: A retains the input arrangement, whereas B swaps the designated AC pair. Besides the residual and low-frequency coefficient features of each candidate, the model receives their corresponding sums ( $A + B$ ) and differences ( $A - B$ ), retaining combined information while exposing differences between arrangements. For residual distributions, the conventional method counts near-integer values at a single threshold. Here, 49 thresholds from 0.01 to 0.49 are used to compare the numbers of residuals below each threshold, so the model can observe distribution differences over multiple ranges. These features are fed to a Transformer, and a shared scorer compares the two candidate orders to predict the swap state.

Carriers are also selected by prediction confidence: the  $K$  most confident blocks are used first, a procedure referred to as Top- $K$  selection. The encoder and decoder identify these blocks using the same rule, without transmitting a complete selection list. For blocks that are still misclassified, the encoder compresses the correction information and stores it in the same marked JPEG. The decoder first reads the message from the coefficient order, then restores the coefficients using the predictions and correction information. We compare classification accuracy and correction-information size, and check whether the restored quantized coefficients and decoded pixels are exactly identical to those before embedding.

## 2. Methods

### 2.1 System scope and coefficient manipulation

The system applies to baseline sequential 4:4:4 color JPEG images with unit quantization tables, in which every quantization value is 1. Embedding modifies only the luminance (Y) channel; the chrominance channels (Cb and Cr) are unchanged. The design is self-contained: the embedded message, control packet, and restoration information are all stored in the same marked JPEG, from which the receiver can extract the message and restore the original image. The decoder still requires the same shared key, trained and deployed model, and decoding program as the encoder.

Carriers are restricted to the two AC coefficients at zigzag indices 1 and 4 in each  $8 \times 8$  DCT block, denoted by  $q_a$  and  $q_b$ . Their relative order represents the embedded bit:  $q_a < q_b$  represents 1, and  $q_a > q_b$  represents 0. If the original order does not represent the desired bit, the coefficients are swapped. To balance image quality and stability, carriers must satisfy the Cap50 condition, with an absolute difference between 4 and 50:

$$b(q) = \mathbf{1}[q_a < q_b], \quad 4 \leq |q_a - q_b| \leq 50. \quad (1)$$


Let  $T$  denote the operation that swaps the two designated coefficients in a block and leaves all other coefficients unchanged. Superscript 0 denotes no swap, and superscript 1 denotes one swap.

Swapping does not change the absolute coefficient difference, i.e.,  $|q'_a - q'_b| = |q_a - q_b|$ . The decoder can therefore identify the same eligible carriers without additional information. For a carrier block  $q$  and message bit  $c \in \{0, 1\}$ , the swap state  $s$  and marked block  $q'$  are defined as

$$s = b(q) \oplus c, \quad q' = T^s(q). \quad (2)$$

The entire procedure operates directly on quantized DCT coefficients (Figure 1). During decoding, the receiver reads the embedded message directly from the relative order of the marked coefficients and restores the original coefficients using predictions and correction information. Here,  $\hat{s}$  is the predicted swap state, and  $e$  is the correction bit that resolves the difference between the predicted and actual states:

$$s = \hat{s} \oplus e, \quad \hat{q} = T^{\hat{s} \oplus e}(q') = q. \quad (3)$$



Encoder and decoder operate on quantized JPEG coefficients; the original JPEG is not provided to the decoder.

Figure 1. Self-contained encoding and decoding. The encoder stores the message and restoration information in the same marked JPEG. The decoder reads the message from the coefficient order, then restores the original coefficients using swap-state predictions and the correction map.

## 2.2 Conventional baseline

A spatial integer-rounding heuristic is used as an external baseline comparator. Given a quantized DCT block  $A$  and the candidate  $B = T(A)$  obtained by swapping AC coefficients (1,4), a standard orthonormal float32 inverse discrete cosine transform (IDCT) produces the corresponding spatial blocks  $I_A$  and  $I_B$ . The rounding residual at pixel  $(m, n)$  is defined as

$$r[m, n] = |I[m, n] - \text{round}(I[m, n])| \in [0, 0.5]. \quad (4)$$

With the integer-proximity threshold set to  $T_{\text{hr}} = 0.3$ , the number of qualifying pixels among the 64 pixels is counted:

$$i_A = \sum_{m=0}^7 \sum_{n=0}^7 \mathbf{1}[r_A[m, n] < T_{\text{hr}}], \quad i_B = \sum_{m=0}^7 \sum_{n=0}^7 \mathbf{1}[r_B[m, n] < T_{\text{hr}}]. \quad (5)$$

Unmodified natural-image blocks have more spatial pixel values close to integers. The predicted swap state is therefore  $\hat{s} = \mathbf{1}[i_B > i_A]$ : if  $i_B > i_A$ , input block  $A$  is classified as swapped ( $\hat{s} = 1$ ); otherwise, if  $i_A \geq i_B$ , it is classified as unswapped ( $\hat{s} = 0$ ).

## 2.3 ARPT architecture and feature construction

ARPT retains the input block as candidate A and swaps its two designated coefficients to form candidate B. It compares these arrangements to determine whether the input block was swapped. The model uses three feature types, comprising 129 feature tokens. Adding one [CLS] token gives an input sequence of length 130 (Figure 2).

**64 spatial tokens:** At each spatial location, the rounding residuals of both candidates, their difference, and their sum describe the local residual relationship:

$$x_{\text{sp},k} = [2r_{A,k}, 2r_{B,k}, 2(r_{A,k} - r_{B,k}), r_{A,k} + r_{B,k}]^\top \in \mathbb{R}^4. \quad (6)$$

Spatial indices range from 0 to 63. The first three components are multiplied by 2 so that the residual components lie in  $[0, 1]$  and their difference in  $[-1, 1]$ . The sum already lies in  $[0, 1]$  and retains its original scale.

**49 distribution-profile tokens:** Spatial tokens retain location-specific residual information, whereas profile tokens compare the residual distributions of the entire block. Thresholds are set at intervals of 0.01 from 0.01 to 0.49. For each threshold, the difference between the numbers of residuals strictly below that threshold is scaled by a threshold-specific standard deviation estimated from the training data. The mean of the difference is fixed at zero, so swapping the candidates negates the distribution-difference vector:

$$\begin{aligned} t_j &= 0.01j, & j &= 1, \dots, 49, \\ d_j &= \sum_{k=0}^{63} \mathbf{1}[r_{A,k} < t_j] - \sum_{k=0}^{63} \mathbf{1}[r_{B,k} < t_j], \\ z_j &= d_j / \sigma_j, & x_{\text{pr},j} &= [z_j, (t_j - 0.25) / 0.24]^\top \in \mathbb{R}^2. \end{aligned} \quad (7)$$

The first component is the scaled residual-count difference. The second maps the threshold coordinate to  $[-1, 1]$ , identifying the threshold associated with each token.

**16 coefficient-context tokens:** Zero-based zigzag indices 0 through 15, including DC, supply low-frequency coefficient information. Both candidates are standardized using shared position-specific means and standard deviations estimated from the training data. The standardized coefficients, their difference, and their sum form the features at each position:

$$x_{cf,m} = [c_{A,m}, c_{B,m}, c_{A,m} - c_{B,m}, c_{A,m} + c_{B,m}]^\top \in \mathbb{R}^4, \quad m = 0, \dots, 15. \quad (8)$$

Here,  $c_{A,m}$  and  $c_{B,m}$  are the standardized coefficients of the two candidates at zigzag position  $m$ .


Each feature type is linearly projected to 64 dimensions. Learnable position and type embeddings are added, and the tokens are combined with the [CLS] token to form the input sequence. The Transformer encoder [5] has 2 layers, 4 attention heads, and a 128-dimensional feed-forward network, with GELU activation and dropout set to 0. The scoring pipeline is

**Feature embedding  $\rightarrow$  Transformer encoder  $\rightarrow$  [CLS] output  $\rightarrow$  scoring head  $\rightarrow$  score**

The shared scorer comprises this entire pipeline. Using the same parameters, the model scores the original candidate order and the reversed order, divides their score difference by 2, and applies a sigmoid to obtain the predicted swap-state probability:

$$f(A, B, z) = \frac{g(A, B, z) - g(B, A, -z)}{2}, \quad p = \text{sigmoid}(f) = \frac{1}{1 + \exp(-f)}. \quad (9)$$

Here,  $g$  is the shared scorer,  $z$  is the distribution-difference vector in Eq. (7), and  $p$  is the probability that the model classifies input block A as swapped. Reversing the candidate order exchanges their features and negates the differences, while leaving the sums unchanged. Consequently, the score difference  $f$  changes sign, giving the model its antisymmetric structure. In direct classification, a probability of at least 0.5 indicates a swapped block. At  $f = 0$ , the probability is exactly 0.5 in either order; the threshold rule therefore does not imply complementary hard labels at this tie.



A and B are two candidate arrangements of the same block; when A and B are swapped, the score changes to its negative.

Figure 2. ARPT architecture. Three feature types are evaluated by a shared model, and the difference between scores for the two candidate orders produces the swap-state prediction.

ARPT has 78,657 trainable parameters and is trained with binary cross-entropy loss without class weighting. The deployed model uses seed 2025, the AdamW optimizer [6], a learning rate of  $3 \times 10^{-4}$ , weight decay of  $10^{-4}$ , and a batch size of 256, for 10 epochs. The model and gating threshold are selected according to the number of post-gating errors on the validation set. The deployed checkpoint is epoch 10, with a threshold of 0.555.

A decision gate combines the conventional method with ARPT. The ARPT decision replaces the conventional decision only when its probability in support of the opposite decision reaches the specified threshold; otherwise, the conventional decision is retained. The conventional decision is used only at this decision stage, not as an ARPT input.

An ARPT output closer to 0 or 1 indicates greater confidence. Its distance from 0.5 is used as the confidence score for carrier selection in the next subsection.

## 2.4 Confidence-guided carrier selection

Blocks for which the model is more confident are used first to reduce the correction information required after misclassification. An output  $p$  close to 0 or 1 indicates a stronger preference for one decision, whereas an output near 0.5 indicates uncertainty. The carrier-selection confidence score is therefore  $|p - 0.5|$ .

To compute confidence consistently before and after embedding, both encoder and decoder first arrange the two designated coefficients of every eligible block in the same order: the smaller coefficient at position 1 and the larger at position 4. The resulting block is then passed to ARPT. Even if embedding swaps the coefficients, this canonical arrangement produces the same model input. To reduce the effect of small floating-point differences on ranking, confidence is converted to an integer:

$$\rho = \lfloor 256|p - 0.5| \rfloor, \quad \rho \in \{0, \dots, 128\}. \quad (10)$$


Eligible blocks are ranked in descending order of model confidence, and the first  $K$  are selected for message embedding (Top-K selection). Ties are resolved by an order generated from the shared key. After selection, the message is embedded in key order: ascending coefficient order represents 1, and descending order represents 0.

The decoder repeats the same selection procedure. To check that the selected blocks and their order agree, the encoder stores the first 16 bytes of the carrier-list SHA-256 digest in the control packet. This digest is used only to check selection agreement. The decoder recomputes and compares it, and stops decoding if it does not match.

## 2.5 Self-contained protocol, restoration, and cost

The Y blocks are partitioned by the shared key into three disjoint regions (Figure 3):

- **Bootstrap region (25%):** Stores a fixed 32-byte bootstrap packet specifying how much control information the decoder should read and supplying a CRC-32 checksum.
- **Control region (50%):** Stores the correction map, message length, correction-map encoding, and the settings and identifiers needed for decoding.
- **Data region (25%):** Embeds the message through coefficient swapping at fixed positions, here zigzag positions 1 and 4.



The three pools are disjoint; Bootstrap and Control use independent reversible AC histogram shifting.

Figure 3. Three disjoint pools and confidence ranking. The bootstrap and control pools carry information through reversible AC histogram shifting; the data pool carries the message through coefficient swapping.

The bootstrap and control regions store information using reversible histogram shifting [7]. Zero coefficients are unchanged. For coefficients equal to 1 or  $-1$ , the magnitude represents one message bit; other coefficients are shifted outward by one to create space:

$$a' = \begin{cases} 0, & a = 0, \\ \text{sign}(a)(1 + d), & |a| = 1, \\ a + \text{sign}(a), & |a| > 1. \end{cases} \quad (11)$$

Here,  $a$  is the original coefficient and  $d$  is the bit, 0 or 1, to be written. For 0, the magnitude stays at 1; for 1, it becomes 2, with the sign unchanged. After extraction, the decoder restores coefficients of magnitude 1 or 2 to magnitude 1 and shifts the other modified coefficients back to their original values.

As shown in Figure 1, the decoder reads the bootstrap packet, reads the control information, identifies the data carriers, extracts the message, and finally restores the original coefficients using ARPT and the correction map. Each region is read according to the length recorded in the packet. Decoding stops if the length or format is inconsistent.

The correction map records which carriers have incorrectly predicted swap states: 0 denotes a correct prediction, and 1 an error. If  $w$  of the  $K$  selected carriers require correction, the correction rate is  $w/K$ .

Four lossless storage methods are compared: direct bit packing (packed), zlib compression, storage of error positions only (sparse), and combinatorial encoding. The encoder chooses the shortest packet and records its encoding in the control region so that the decoder can reconstruct the map.

Combinatorial encoding [8] records the index of the error-position set among all possible sets of  $w$  errors in  $K$  positions. For error positions  $0 \leq a_1 < \dots < a_w < K$ , the index and required bit count are

$$\text{rank} = \sum_{i=1}^w \binom{a_i}{i}, \quad 0 \leq \text{rank} < \binom{K}{w}, \quad B_{\text{rank}} = \left\lceil \log_2 \binom{K}{w} \right\rceil. \quad (12)$$

The error count  $w$  is stored separately as an unsigned LEB128 integer. With  $K$ ,  $w$ , and the combinatorial rank, the decoder reconstructs the error positions.

To compare overhead at the same embedded message size, overhead-adjusted payload is defined as the actual message size  $K$  minus the sizes of the control and bootstrap packets:

$$P_{\text{adj}} = K - B_{\text{control}} - B_{\text{bootstrap}}, \quad B_{\text{bootstrap}} = 256 \text{ bits.} \quad (13)$$

The message actually written and extracted is still the  $K$  bits in the data region. Because the bootstrap, control, and data regions are physically disjoint pools, overhead-adjusted payload is only an accounting metric; its increase does not constitute experimental evidence of an increase in the maximum embeddable message size. A larger  $P_{\text{adj}}$  means less packet overhead for the same message size. The control-packet size already includes the correction map, which is not subtracted again.

Restoration is checked using the JPEG files actually stored. The extracted message must match the original message, and the restored Y, Cb, and Cr quantized coefficients and RGB pixels produced by a fixed decoder must each be identical to their pre-embedding counterparts. Exact restoration is declared only when all three checks pass.

## 2.6 Data and comparison design

The study uses 37 iPhone 14 Pro images, with 2,000 coefficient pairs constructed per image, for 74,000 pairs. Each pair has an unswapped and a swapped state, yielding 148,000 samples labeled 0 and 1, respectively.

The data are split by image into 25 training images, 6 validation images, and 6 test images, with no overlap between the three sets.

The main comparison fixes the embedded message at 2,560 bytes, or  $K = 20,480$  bits, and evaluates three methods:

- **Conventional method:** Selects blocks in a key-generated pseudorandom order and predicts swap states using the conventional comparator.
- **Keyed ARPT:** Selects blocks in the same key-generated pseudorandom order and predicts swap states using gated ARPT.
- **Confidence-guided ARPT:** Uses the same gated ARPT predictor but prioritizes blocks with higher model confidence.

The first two methods isolate the difference between predictors; the latter two isolate the effect of confidence ranking. All three use the same correction-map encoding selection procedure, and the cost of each method’s control packet is included. A separate comparison between the conventional method and gated ARPT uses an earlier encoding configuration; its results are reported separately.

A further 32 images from the same device are used for capacity analysis, and 24 JPEG images and 100 DIV2K images [9] are used for stress testing. None of these images is used for training or validation. All are converted to JPEG with the same unit quantization tables and 4:4:4 format.

The original target for each of the 32 images was 2,560 bytes. Some images had too few eligible blocks, so the message size was subsequently adjusted to the available block count, with a maximum of 2,560 bytes:

$$L_i = \min(2560, \lfloor N_i/8 \rfloor) \text{ bytes}, \quad K_i = 8L_i. \quad (14)$$



Here,  $N_i$  is the number of eligible blocks in the data region of image  $i$ , and  $L_i$  is the embeddable byte count. This adjustment was made after the original capacity failures, and the results are reported as a post hoc analysis.

Per-image gain is the difference in overhead-adjusted payload between ARPT and the conventional method. Paired bootstrap resampling at the image level estimates the 95% confidence interval for the mean gain. For the 24 images that succeeded at the fixed message size, the original results from 10,000 resamples with random seed 2025 are retained. For the 32 capacity-adjusted images, 100,000 resamples are used with random seed 20260918.

The correction-map encoding experiment uses the six validation images. The same map from each image is encoded by all four methods and compared with automatic selection, recording the required bit counts.

### 3. Experimental results

#### 3.1 Marked images and swap-state classification

Figure 4 shows three watermarked JPEG images. Confidence-guided ARPT selects the blocks in each image. The left and right images each carry 2,560 bytes, and the middle image carries 2,448 bytes.



Marked JPEGs | Confidence-guided ARPT | Left/right: 2,560 bytes each; middle: 2,448 bytes

Figure 4. Three watermarked JPEG images. The left and right images each carry a 2,560-byte message; the middle image carries a 2,448-byte message.

Swap-state classification is evaluated on 12,000 coefficient pairs from the six test images. Both the unswapped and swapped states are tested for every pair, giving 24,000 decisions.

The conventional method achieves 91.17% accuracy, with 2,119 errors. Direct ARPT classification achieves a mean accuracy of 94.66% over three training seeds, with a mean of 1,282 errors, 837 fewer than the conventional method.

Gated ARPT combines the conventional decision with the model prediction: it retains the conventional decision unless ARPT’s confidence in the opposite decision reaches the specified threshold. Over the three seeds, its mean accuracy is 94.62%, with a mean of 1,291 errors.

Direct ARPT outperforms the conventional method for all three seeds, but gating does not consistently improve the result. Subsequent carrier-selection and encoding–decoding experiments use the seed-2025 model with the gating threshold selected on the validation data.

Table 1. Swap-state classification on the test set.

| Seed | Direct ARPT accuracy | Errors | Gated ARPT accuracy | Errors |
|------|----------------------|--------|---------------------|--------|
| 2025 | 94.71%               | 1,270  | 94.76%              | 1,258  |
| 2026 | 94.55%               | 1,308  | 94.53%              | 1,313  |
| 2027 | 94.72%               | 1,268  | 94.58%              | 1,302  |
| Mean | 94.66%               | 1,282  | 94.62%              | 1,291  |

### 3.2 Carrier selection and restoration cost on the test set

The conventional method, Keyed ARPT, and confidence-guided ARPT are compared on the six test images. Each method embeds 2,560 bytes per image, with  $K = 20,480$  bits fixed.

The conventional method and Keyed ARPT select the same blocks, with selected-carrier error rates of 8.77% and 5.45%, respectively. With the same gated ARPT predictor, confidence-guided selection further reduces the error rate to 0.66% (Figure 5). The mean numbers of blocks requiring correction per image are 1,795.17, 1,116.33, and 134.50, respectively.




Figure 5. Carrier-selection results and restoration-information costs for the six test images. All three methods embed 20,480 bits per image. Values are means over the six images.

As Table 2 shows, the mean correction-map size falls from 6,468.00 bits for Keyed ARPT to 1,345.33 bits for confidence-guided ARPT, a reduction of 5,122.67 bits per image. Confidence-guided selection additionally stores a 128-bit carrier-list digest. Even with this cost included, the control packet decreases from 6,772.00 to 1,777.33 bits.

Table 2. Mean restoration-information costs and exact-restoration results for the six test images.

| Metric                           | Conventional method | Keyed ARPT | Confidence-guided ARPT |
|----------------------------------|---------------------|------------|------------------------|
| Correction-map size (bits)       | 8,996.00            | 6,468.00   | 1,345.33               |
| Control-packet size (bits)       | 9,172.00            | 6,772.00   | 1,777.33               |
| Overhead-adjusted payload (bits) | 11,052.00           | 13,452.00  | 18,446.67              |
| Exact-restoration cases          | 6/6                 | 6/6        | 6/6                    |

All three methods embed 20,480 bits per image. The bootstrap packet is 256 bits, and the control packet includes the correction map.

Relative to Keyed ARPT, confidence-guided selection increases mean overhead-adjusted payload by 4,994.67 bits. All 18 encoding-decoding cases for the three methods recover the original message and restore exactly the same Y, Cb, and Cr quantized coefficients and decoded RGB pixels as before embedding.

### 3.3 Correction-map encoding

The same correction map from each of the six validation images is encoded by different methods to compare storage requirements.

Table 3 shows that combinatorial encoding has the smallest mean correction-map size, 1,938.67 bits.

Table 3. Mean costs of correction-map encodings on the six validation images.

| Encoding                         | Correction-map size<br>(bits) | Control-packet size (bits) | Overhead-adjusted<br>payload (bits) |
|----------------------------------|-------------------------------|----------------------------|-------------------------------------|
| packed                           | 20,688.00                     | 21,120.00                  | −896.00                             |
| zlib                             | 2,926.67                      | 3,358.67                   | 16,865.33                           |
| sparse                           | 2,373.33                      | 2,805.33                   | 17,418.67                           |
| combinatorial                    | 1,938.67                      | 2,370.67                   | 17,853.33                           |
| Automatic encoding<br>(new-auto) | 1,938.67                      | 2,370.67                   | 17,853.33                           |

Automatic selection (new-auto) compares packet sizes and chooses the shortest encoding. All six images select combinatorial encoding in this experiment, so the results are identical. The total overhead of packed encoding exceeds the embedded message size, and it is not selected as the final encoding.

Five methods are tested on each of the six images, for 30 encoding–decoding cases. All recover the message and exactly restore the original quantized coefficients and decoded RGB pixels.

### 3.4 Fixed message size and capacity adjustment

The conventional method and gated ARPT are compared on a further 32 JPEG images from the same device. Automatic correction-map encoding (legacy-auto) selects the shortest encoded packet.

The original target was 2,560 bytes per image, but 8 images had too few eligible blocks to accommodate this message, so the goal of success on every image was not met. For the remaining 24 images, both methods achieve exact restoration in all 48 encoding–decoding cases. Gated ARPT increases mean overhead-adjusted payload by 2,539.33 bits per image, with a 95% confidence interval of 2,329.33–2,777.01 bits.

To handle insufficient block counts, the message size is subsequently adjusted using Eq. (14): images with enough available blocks still carry 2,560 bytes, whereas the others carry the amount they can accommodate. Actual embedded messages range from 2,139 to 2,560 bytes, with a mean of 2,491.69 bytes. Both methods recover the message and exactly restore the image in all 64 cases over the 32 images.

Figure 6 shows the per-image gains after message-size adjustment. Gated ARPT increases mean overhead-adjusted payload by 2,363.25 bits per image, with a 95% confidence interval of 2,168.75–2,576.00 bits. The gain is positive for all 32 images.




Figure 6. Per-image gains of gated ARPT over the conventional method after message-size adjustment. The red line indicates the mean gain of 2,363.25 bits; the shaded band indicates the 95% confidence interval for the mean gain.

### 3.5 Tests on different image content

The conventional method and confidence-guided ARPT are compared on 32 outdoor JPEG images, 24 indoor JPEG images, and 100 DIV2K images.

Both methods use automatic encoding (new-auto), and the message size for each image is set by Eq. (14). Confidence-guided ARPT uses the gated predictor and selects carriers by model confidence.

As Table 4 shows, mean overhead-adjusted payload increases in all three image sets. Mean gains per image are 5,288.50 bits for the 32 outdoor JPEG images, 2,918.00 bits for the 24 indoor JPEG images, and 461.92 bits for DIV2K.

Table 4. Mean overhead-adjusted payload and exact-restoration results for different image sets.

| Image set               | Conventional method<br>(bits/image) | Confidence-guided ARPT (bits/image) | Mean gain<br>(bits/image) | Exact-restoration cases |
|-------------------------|-------------------------------------|-------------------------------------|---------------------------|-------------------------|
| Outdoor JPEG, 32 images | 10,803.00                           | 16,091.50                           | 5,288.50                  | 64/64                   |
| Indoor JPEG, 24 images  | 16,415.67                           | 19,333.67                           | 2,918.00                  | 48/48                   |
| DIV2K, 100 images       | 2,383.28                            | 2,845.20                            | 461.92                    | 200/200                 |

The indoor and outdoor photographs from the same device are predominantly  $4032 \times 3024$ , approximately 12 megapixels. The DIV2K images all have a long side of 2040 pixels and average approximately 2.84 megapixels. Because the DIV2K images are smaller, fewer blocks are available for embedding, and the actual embedded messages are smaller.

All 312 encoding–decoding cases across the three sets recover the message and exactly restore the original quantized coefficients and decoded RGB pixels.

Actual embedded messages in DIV2K range from 107 to 1,227 bytes per image, with a mean of 696.46 bytes. Overhead-adjusted payload increases for 98 of the 100 images; for the other two, it decreases by 16 and 112 bits. A possible explanation is that the correction-map savings do not offset the additional model and carrier-check information.

## 4. Conclusions and future directions

This study proposes ARPT and confidence-guided carrier selection to reduce restoration-information costs in reversible JPEG watermarking. ARPT compares two candidate block arrangements and uses residual and low-frequency coefficient features to classify the swap state. Confidence-guided selection prioritizes blocks for which the model is more confident, reducing the correction information that must be stored. Remaining errors are recorded in a correction map in the same marked JPEG, allowing the decoder to restore the original image.

On the six test images, direct ARPT achieves a mean classification accuracy of 94.66%, compared with 91.17% for the conventional method. With the same predictor and fixed message size, confidence-guided selection reduces the selected-carrier error rate from 5.45% to 0.66% and increases mean overhead-adjusted payload from 13,452.00 to 18,446.67 bits. All 18 encoding-decoding cases for the three methods recover the message and exactly restore the original quantized coefficients and decoded RGB pixels.

Tests on indoor, outdoor, and DIV2K images also show lower mean restoration-information costs. The message size an image can accommodate still depends on its eligible block count, however. At a fixed target of 2,560 bytes, some images fail because there are too few blocks; complete encoding, decoding, and restoration are achieved only after adjusting the message size to the available blocks. We also tested selecting fewer blocks, adding an error-risk selector, and introducing a ranking loss, but none provided sufficient payload benefit to replace the current confidence-guided strategy.

The present experiments use 4:4:4 JPEG images with all quantization values equal to 1 and embed messages by swapping AC coefficients at fixed positions. Future work could evaluate other quantization tables, chroma-sampling schemes, and coefficient positions to assess classification and restoration under different JPEG settings. Message size could also be chosen according to each image’s available blocks and correction-information cost. Implementation work could reduce inference time, test carrier-selection agreement on different hardware, and investigate message verification and protection against transmission errors.

## References

- [1] S. Weng, Y. Zhou, T. Zhang, M. Xiao, and Y. Zhao, “Reversible Data Hiding for JPEG Images With Adaptive Multiple Two-Dimensional Histogram and Mapping Generation,” *IEEE Transactions on Multimedia*, vol. 25, pp. 8738–8752, 2023. DOI: 10.1109/TMM.2023.3241541.
- [2] F. Li, Z. Qi, X. Zhang, and C. Qin, “Progressive Histogram Modification for JPEG Reversible Data Hiding,” *IEEE Transactions on Circuits and Systems for Video Technology*, vol. 34, pp. 1241–1254, 2024. DOI: 10.1109/TCSVT.2023.3288038.
- [3] Z.-H. Ou and L.-H. Chen, “A Robust Watermarking Method for Stereo-Pair Images Based on Unmatched Block Bitmap,” *Multimedia Tools and Applications*, vol. 75, pp. 3259–3280, 2016. DOI: 10.1007/s11042-014-2433-0.
- [4] Z.-H. Ou, “A Novel Reversible Watermarking Method for Stereo-Pair Image Based on Ac-Pair Swapping,” in *2024 International Conference on Machine Learning and Cybernetics (ICMLC)*, pp. 46–51, 2024. DOI: 10.1109/ICMLC63072.2024.10935238.
- [5] A. Vaswani et al., “Attention Is All You Need,” in *Advances in Neural Information Processing Systems 30*, pp. 5998–6008, 2017. NeurIPS proceedings.

- [6] I. Loshchilov and F. Hutter, “Decoupled Weight Decay Regularization,” in *International Conference on Learning Representations (ICLR)*, 2019. arXiv:1711.05101.
- [7] Z. Ni, Y.-Q. Shi, N. Ansari, and W. Su, “Reversible Data Hiding,” *IEEE Transactions on Circuits and Systems for Video Technology*, vol. 16, no. 3, pp. 354–362, 2006. DOI: 10.1109/TCSVT.2006.869964.
- [8] T. M. Cover, “Enumerative Source Encoding,” *IEEE Transactions on Information Theory*, vol. 19, no. 1, pp. 73–77, 1973. DOI: 10.1109/TIT.1973.1054929.
- [9] E. Agustsson and R. Timofte, “NTIRE 2017 Challenge on Single Image Super-Resolution: Dataset and Study,” in *2017 IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pp. 1122–1131, 2017. DOI: 10.1109/CVPRW.2017.150.